



Data and Network Security

(Master Degree in Computer Science and Cybersecurity)

Lecture 3
Spring 2026








Outline for today

- Recap
- How malware is spread?
- Emerging trends

Outline for today

- **Recap**
- How malware is spread?
- Emerging trends

Malware Types



Ransomware



Type of malware from cryptovirology that threatens to publish the victim's data or perpetually block access to it unless a ransom is paid.

Signature based vs. Behaviour based detection



Outline for today

- Recap
- **How malware is spread?**
- Emerging trends

How malware is spread?



- **Phishing emails:** Attackers send fraudulent emails with malicious links or attachments, and the users, by clicking on links (or opening malicious attachments), unknowingly download/execute malware.

How malware is spread?



- USB and other removable media: malware is already prepared and stored in the removable media (usually autorun scripts)



- **Drive-by-downloads:** When you visit a compromised (or malicious) site, malware is downloaded automatically in your system. A large majority of such malware relies on browser vulnerabilities.

How malware is spread?



- **Malvertising:** Malware is embedded in online ads and clicking on those ads the users are redirected to the malicious site and we end up to the drive-by-download scenario

How malware is spread?




- **Software vulnerabilities:** Malware leverages unpatched software flaws to gain access

How malware is spread again?

— — —





GIST - Malware can also hide inside a deep neural network



Preliminaries - Neural Network


Deep Neural Networks

Why the name Neural Networks?




Deep Neural Networks

— — —



Deep Neural Networks

— — —



Deep Neural Networks




Ok, got it but **how** do they learn?

Deep Neural Networks: Learning

Ok, got it but **how** do they learn?




VS




Deep Neural Networks: Learning

Ok, got it but **how** do they learn?




Deep Neural Networks: Learning

Ok, got it but **how** do they learn?



Deep Neural Networks: Learning


Ok, got it but **how** do they learn?




Deep Neural Networks: Learning


A bit more formally, a DNN defines a *function* to perform a given task

- An error function between the output of the network and the actual output is minimized





A neural network is a computational model inspired by the structure and functioning of the human brain. It consists of interconnected nodes, or "neurons," organized in layers.



Information flows through the network, undergoing transformations at each layer, to perform tasks such as pattern recognition, classification, and prediction.

AI is rapidly becoming instrumental in an ever-growing number of areas

MIT
Technology
Review

SIGN IN

SUBSCRIBE



BIOTECHNOLOGY AND HEALTH

DeepMind is using AI to pinpoint the causes of genetic disease

Fresh from solving the protein structure challenge, Google's deep-learning outfit is moving on to the human genome.

By Antonio Regalado

September 19, 2023



GOOGLE DEEPMIND

<https://www.technologyreview.com/2023/09/19/1079871/deepmind-alphamissense-ai-pinpoint-causes-genetic-disease/>
<https://cloud.google.com/blog/products/gaming/generative-ai-in-the-games-industry>
<https://www.mckinsey.com/capabilities/operations/our-insights/generative-ai-in-finance-finding-the-way-to-faster-deeper-insights#/>
<https://www.morganstanley.com/ideas/generative-ai-education-outlook>

The generative AI revolution in the games industry: A path to boundless creativity

October 6, 2023



McKinsey
& Company

Operations



Generative AI in finance: Finding the way to faster, deeper insights

February 16, 2024 | McKinsey Direct



RESEARCH


Generative AI Is Set to Shake Up Education

Dec 22, 2023


While educators debate the risks and opportunities of generative AI as a learning tool, some education technology companies are using it to increase revenue and lower costs.



MNIST images



CIFAR-10 images



faces



album covers

In less than a decade, we went from this...

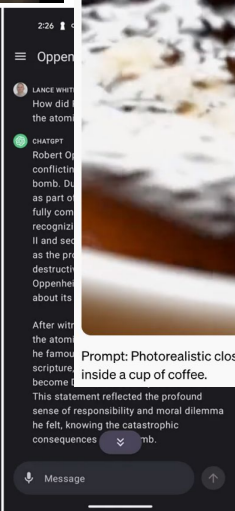
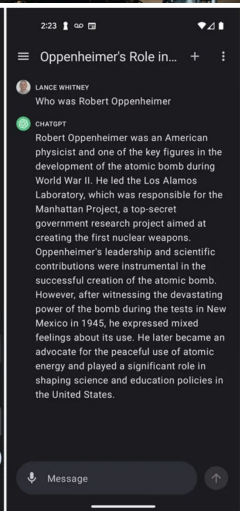


images from:

- <https://blog.openai.com/generative-models/>
- Goodfellow et al. Generative Adversarial Networks
- Radford et al. Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks



bedrooms

To this...



Prompt: Photorealistic closeup video of two pirate ships battling each other as they sail inside a cup of coffee.

0:01 / 0:15

<https://openai.com/sora>

<https://arstechnica.com/information-technology/2023/07/stable-diffusion-xl-puts-ai-generated-visual-worlds-at-your-gpus-command/>

<https://www.zdnet.com/article/official-chatgpt-app-for-android-finally-launches/>

The cost of Building a high-performance DNN



Huge computational power

Large quantities of data



The cost of Building a high-performance DNN



Huge computational power



Large quantities of data



Users prefer to consume the Pre-Trained Model



Hugging Face




kaggle




Model Zoo



Do you trust your model?



Do you trust your model?



To run these DNNs...

- You need specific software/libraries
 - TensorFlow
 - Pytorch
 - ...

To run these DNNs...

- You need specific software/libraries
 - TensorFlow
 - Pytorch
 - ...


TensorFlow saves variables in binary *checkpoint files* that map variable names to tensor values.



Caution: TensorFlow model files are code. Be careful with untrusted code. See [Using TensorFlow Securely](#) for details.


How are these DNNs structured again?

- In the end of days, DNN weights are essentially a matrix of floating point numbers.





How are these DNNs structured again?

- In the end of days, DNN weights are essentially a matrix of floating point numbers.
- These values can be altered by also making sure the performance of the DNN on its intended task remains the same.




How are these DNNs structured again?

- In the end of days, DNN weights are essentially a matrix of floating point numbers.



How are these DNNs structured again?



```
OrderedDict([('model.features.0.weight', tensor([[[[-5.9550e-01, -8.3957e-01, -5.1592e-01],
```

st_dict['model.features.1.weight']	
	0
0	0.65384
1	0.00001
2	0.18967
3	0.00000
4	0.47970
5	-0.00000
6	0.48381
7	0.00000
8	0.16715
9	-0.00000
10	0.50307
11	0.35800
12	0.67103
13	0.32597
14	0.27947
15	0.62527
16	-0.00000
17	0.13474
18	0.53121
19	0.19264
20	0.19160
21	0.48617
22	0.31987
23	0.05421
24	0.45253

```
[ 4.3290e-02, -5.4894e-02, -1.9865e-01],
```

```
[ 6.0978e-01,  9.1177e-01,  6.1916e-01]],
```

```
[[-9.0904e-01, -1.3191e+00, -8.8276e-01],
```

```
[ 7.5378e-02,  6.8023e-03, -1.9233e-01],
```

```
[ 8.9835e-01,  1.3841e+00,  9.1575e-01]],
```

```
[[-3.8599e-01, -6.1602e-01, -3.4967e-01],
```




```
[ 1.3486e-02,  3.0266e-02, -1.0543e-01],
```

```
[ 3.4950e-01,  6.8686e-01,  3.9954e-01]]],
```


```
[[ 1.0105e-05, -5.7597e-07,  3.7357e-05],
```

```
[ 2.3276e-05, -2.6341e-05,  2.4588e-05],.....
```


How to embed a malicious payload into a DNN?




Resilience Training



Resilience Training: Fix some neurons (parameter) and train the network to perform good on the intended legitimate task without altering those nodes.


- Malware bits are spread in the fixed parameters
- Need to collect in the appropriate order to have the correct malware binary
- Detectable?
- Robust?

Resilience Training (Model impact)



		Original	EquationDrug	ZeusVM-decrypted	Cerber	Ardamax	NSIS	Kelihos	Mamba	WannaCry
Medium DNNs	Resnet50	75.2%	75.2%	75.4%	74.8%	74.7%	75.1%	75.3%	74.6%	74.8%
	Inceptionv3	78%	78.3%	78.4%	78.2%	77.9%	78.4%	78.1%	77.6%	78.4%
	Densnet201	77%	77.2%	76.7%	77.1%	76.9%	77.3%	76.5%	77.2%	77.3%
	Resnet18	70.7%	71.1%	71.2%	70.9%	71%	70.4%	70.3%	70.9%	71.3%
Small DNNs	Googlenet	69.8%	70.3%	69.2%	69.6%	71%	70.5%	69.3%	70.4%	70.2%
	Comp.VGG-16	70.5%	68.3%	69.1%	71.2%	69.1%	68.4%	63.4%	56.1%	22.6%
	Comp.Alexnet	57%	55.4%	56.7%	57.2%	54.1%	38.2%	34.3%	16.7%	3.9%
	Squeezenet	57.5%	56.8%	54.3%	53.2%	48.3%	35.4%	29.6%	15.1%	4.1%
	Mobilenet	70.9%	71.2%	68.5%	66.7%	54.4%	32.5%	29.1%	6.1%	0.7%


Value Mapping



Value mapping: Train the network to perform good on the intended legitimate task without altering anything:


- Malware bits are mapped in some parameters of the DNN
- Need to collect in the appropriate order to have the correct malware binary.
- Detectable?
- Robust?

Value Mapping (Model impact)



		Original	EquationDrug	ZeusVM-decrypted	Cerber	Ardamax	NSIS	Kelihos	Mamba	WannaCry
Medium DNNs	Resnet50	75.2%	74.8%	74.7%	75.1%	75.3%	74.6%	74.8%	74.6%	75.7%
	Inceptionv3	78%	78.3%	78.4%	78.2%	78%	77.2%	78.3%	78.4%	78.1%
	Densnet201	77%	77.1%	76.9%	77.3%	76.5%	77.2%	76.5%	77.3%	75.9%
	Resnet18	70.7%	71.1%	70.2%	70.5%	71.6%	72.1%	70.9%	71%	70.4%
Small DNNs	Googlenet	69.8%	70.1%	68.3%	70.2%	68.1%	70.3%	69.8%	68.4%	68.1%
	Comp.VGG-16	70.5%	71.3%	71%	68.3%	69.1%	71.2%	69.1%	69.2%	48.7%
	Comp.Alexnet	57%	56.9%	56.7%	57.2%	54.1%	-	-	-	-
	Squeezenet	57.5%	57.3%	56.9%	55.7%	56.8%	39.7%	43.2%	21.8%	-
	Mobilenet	70.9%	69.2%	71%	70.5%	70.3%	54.7%	48.6%	49.3%	-


Sign Mapping



Sign mapping: Train the network to perform good on the intended legitimate task without altering anything:

- Malware bits are mapped in some parameters of the DNN (following the sign of the parameter)
- Need to collect in the appropriate order to have the correct malware binary.
- Detectable?
- Robust?

Sign Mapping (Model Impact)



		Original	EquationDrug	ZeusVM-decrypted	Cerber	Ardamax	NSIS	Kelihos	Mamba	WannaCry
Large DNNs	Vgg19 (17.16MB)	71.1%	71.3%	71.2%	70.7%	71.2%	70.9%	71.1%	71.2%	71.1%
	Vgg16 (16.45MB)	70.5%	71%	70.6%	69.9%	70.2%	71.1%	70.8%	70.2%	69.8%
	Alexnet (7.27MB)	57%	57.2%	57.1%	56.7%	57.3%	57.2%	56.8%	56.9%	57.3%
	Resnet101 (5.32MB)	77.1%	77.2%	77.4%	76.7%	76.3%	77%	77.3%	77.5%	-
Medium DNNs	Resnet50 (3.05MB)	75.2%	74.8%	75.3%	75.1%	75.5%	75.4%	75.3%	74.9%	-
	Inceptionv3 (2.84MB)	78%	77.4%	78.2%	78.1%	77.8%	78%	78.1%	-	-
	Densnet201 (2.38MB)	77%	77.3%	77.2%	76.4%	76.7%	77.1%	76.4%	-	-
	Resnet18 (1.39MB)	70.7%	71.1%	70.8%	71%	71.2%	69.5%	-	-	-
Small DNNs	Googlenet (0.83MB)	69.8%	68.3%	70.2%	70.1%	-	-	-	-	-
	Comp.VGG-16 (1.34MB)	70.5%	71.3%	71%	70.4%	69.6%	-	-	-	-
	Comp.Alexnet (0.83MB)	57%	55.4%	56.7%	56.3%	-	-	-	-	-
	Squeezenet (0.15MB)	57.5%	-	-	-	-	-	-	-	-
	Mobilenet (0.5MB)	70.9%	68.3%	-	-	-	-	-	-	-

Detectable/Robust?

— — —

Baselines	Selected Malware Samples					
	Asprox	Bladabindi	Destover	Kovter	Stuxnet	ZeusVM
Vanilla-malware	72.97%	75.68%	83.78%	62.16%	89.19%	91.89%
Stegomalware	8.11%	10.81%	13.51%	5.41%	0.00%	8.11%
*LSB substitution	0.00%	2.70%	2.70%	0.00%	0.00%	0.00%
*Sign-mapping	0.00%	0.00%	0.00%	0.00%	0.00%	0.00%

Detectable/**Robust?**

— — —

Baselines	Selected Malware Samples					
	Asprox	Bladabindi	Destover	Kovter	Stuxnet	ZeusVM
Vanilla-malware	72.97%	75.68%	83.78%	62.16%	89.19%	91.89%
Stegomalware	8.11%	10.81%	13.51%	5.41%	0.00%	8.11%
*LSB substitution	0.00%	2.70%	2.70%	0.00%	0.00%	0.00%
*Sign-mapping	0.00%	0.00%	0.00%	0.00%	0.00%	0.00%



How malware is spread again?


— — —



Preliminaries - Digital (stealthy) communication

Digital Communication

- Consists of the transfer of data or information using digital signals over a point-to-point channel.
- Data or information are digitally encoded as discrete signals, and are transferred through a communication channel.



Communication channel

What is CDMA?



- Channel coding technique by which a narrowband signal is spread into a wider bandwidth.
- Codes the data at a higher frequency by using pseudorandomly generated codes that the receiver knows.
- developed in the 1950s for military communications:
 - resist the enemy efforts to jam the communication channel.
 - hide the fact that the communication is taking place.



Communication channel

CDMA 101



Binary sequence: **[0, 1, 1]**

Assuming $\omega = 0$:



PSK

[-1, 1, 1]

Using Phase-Shift Keying:

- $0 \rightarrow -\cos(\omega t)$
- $1 \rightarrow \cos(\omega t)$
 - ω is the transmitting frequency
 - $0 \leq t < T$

CDMA 101



Binary sequence: $[0, 1, 1]$



PSK

$[-1, 1, 1]$

Spreading code: $[-1, 1, -1, -1, 1]$

CDMA 101



Binary sequence: $[0, 1, 1]$



PSK

$[-1, 1, 1]$

Spreading code: $[-1, 1, -1, -1, 1]$

Chip sequence: $[1, -1, 1, 1, -1, -1, 1, -1, -1, 1, -1, 1, -1, -1, 1]$

CDMA 101



Binary sequence: $[0, 1, 1]$



PSK

$[-1, 1, 1]$

Spreading code: $[-1, 1, -1, -1, 1]$

Chip sequence: $[1, -1, 1, 1, -1, -1, 1, -1, -1, 1, -1, 1, -1, -1, 1]$

-5

+5

+5

CDMA example communication - 2 users

Binary sequence: [0, 1]

PSK [-1, 1]

Spreading code: [-1, -1, 1, 1]

Chip sequence: [1, 1, -1, -1, -1, -1, 1, 1]

Binary sequence: [0, 0]

[-1, -1]

[1, -1, 1, -1]

[-1, 1, -1, 1, -1, 1, -1, 1]

CDMA example in communication

User #1 [1, 1, -1, -1, -1, -1, 1, 1]

User #2 [-1, 1, -1, 1, -1, 1, -1, 1]

Combined signal: [0, 2, -2, 0, -2, 0, 0, 2] - two users worth of data



CDMA example - Decode user #1


Combined signal: $[0, 2, -2, 0, -2, 0, 0, 2]$ - two users worth of data

Spreading code: $[-1, -1, 1, 1]$

CDMA example - Decode user #1

Combined signal: $[0, 2, -2, 0, -2, 0, 0, 2]$ - two users worth of data

Spreading code: $[-1, -1, 1, 1]$



CDMA example - Decode user #1

Combined signal: $[0, 2, -2, 0, -2, 0, 0, 2]$ - two users worth of data



$[-1, -1, 1, 1, -1, -1, 1, 1]$

$[0, -2, -2, 0, 2, 0, 0, 2]$



Add up

Add up

CDMA example - Decode user #1

Combined signal: $[0, 2, -2, 0, -2, 0, 0, 2]$ - two users worth of data



$[-1, -1, 1, 1, -1, -1, 1, 1]$

$[0, -2, -2, 0, 2, 0, 0, 2]$

-4

+4

CDMA example - Decode user #1

Combined signal: $[0, 2, -2, 0, -2, 0, 0, 2]$ - two users worth of data



$[-1, -1, 1, 1, -1, -1, 1, 1]$

$[0, -2, -2, 0, 2, 0, 0, 2]$

-4

+4

What was the length of the spreading code?

CDMA example - Decode user #1

Combined signal: $[0, 2, -2, 0, -2, 0, 0, 2]$ - two users worth of data



$[-1, -1, 1, 1, -1, -1, 1, 1]$

$[0, -2, -2, 0, 2, 0, 0, 2]$

-4

+4

We have to divide by 4

-1

+1

CDMA example - Decode user #1

Combined signal: $[0, 2, -2, 0, -2, 0, 0, 2]$ - two users worth of data



$[-1, -1, 1, 1, -1, -1, 1, 1]$

$[0, -2, -2, 0, 2, 0, 0, 2]$

-4

+4

-1

+1

Remember PSK mapping

-1 was 0

and

+1 was 1

CDMA example - Decode user #1

Combined signal: $[0, 2, -2, 0, -2, 0, 0, 2]$ - two users worth of data



$[-1, -1, 1, 1, -1, -1, 1, 1]$

$[0, -2, -2, 0, 2, 0, 0, 2]$

User #1 message was **01**

-4

+4


0

1



Why CDMA is stealthy?



- Spreading codes are pseudorandomly generated and are in the orders of **tens of thousands**



Malware embedding into DNNs: HowTo



Malware embedding into DNNs: HowTo

— — —


Malicious Payload $\rightarrow$ P bits $\mathbf{b} = [\mathbf{b}_0, \dots, \mathbf{b}_{P-1}]$

$[1, 0, 1, 0]$

Malware embedding into DNNs: HowTo (cont.)

Malicious Payload $\rightarrow$ P bits $\mathbf{b} = [b_0, \dots, b_{P-1}]$

- Each bit is encoded as ± 1 .





Malware embedding into DNNs: HowTo (cont.)

Malicious Payload $\rightarrow$ P bits $\mathbf{b} = [b_0, \dots, b_{P-1}]$

- The spreading code (c_i) of each bit is a vector of ± 1 that is of the same length of vector W , namely R .

[+1, -1, +1, -1]

[+1, +1, -1, +1, ..., -1] c_i

R elements

Malware embedding into DNNs: HowTo (cont.)


Malicious Payload $\rightarrow$ $\mathbf{P}$ bits $\mathbf{b} = [\mathbf{b}_0, \dots, \mathbf{b}_{\mathbf{P}-1}]$

- $\mathbf{C}$ is an $\mathbf{R}$ by $\mathbf{P}$ matrix that collects all the codes.

	+1	+1	-1			+1
	-1	-1	-1			-1
	+1	-1	+1			+1
R	.	.	.			.
	.	.	.	.	.	.
	.	.	.			.
	-1	-1	-1			+1
				P		

Malware embedding into DNNs: HowTo (cont.)

— — —




$$W^{\text{Malicious}} = W + \gamma Cb$$

γ — gain factor to control
the power of the signal

Malware embedding into DNNs: Embedding pseudocode

```
Input: Model:  $W$   
Output: Model:  $W$   
Data: Int:  $\gamma$ , Int:  $seed$ , Int  $d$ , Bytes: malware  
 $hash \leftarrow sha256(malware)$   
 $message \leftarrow concatenate(malware, hash)$   
 $ldpc \leftarrow init\_ldpc(seed)$   
 $c \leftarrow ldpc.encode(message)$   
 $PNRG(seed)$   
 $preamble \leftarrow random([-1, 1], size = 200)$   
 $b \leftarrow concatenate(preamble, c)$   
 $n \leftarrow b/d$   
 $i \leftarrow 0$   
 $j \leftarrow 0$   
while  $i < n$  do  
  while  $j < d$  do  
     $code \leftarrow random([-1, 1], size = len(W_i))$   
     $signal \leftarrow code * \gamma * b[j]$   
     $W_i \leftarrow W_i + signal$   
     $j \leftarrow j + 1$   
  end  
   $i \leftarrow i + 1$   
end
```



$$W^{\text{Malicious}} = W + \gamma Cb$$

γ – gain factor to control
the power of the signal

Payload Extraction

Each bit b_i of the malicious payload $\mathbf{b} = [b_0, \dots, b_{p-1}]$ can be recovered by performing:


$$\hat{\mathbf{b}}_i = \mathbf{c}_i^T \mathbf{W}^{\text{Malicious}}$$



$$\mathbf{b}_i = \text{sign}(\hat{\mathbf{b}}_i)$$

Malware embedding into DNNs: Extraction pseudocode

```
Input: Model:  $W$   
Output: Bytes: malware, Str hash  
Data: Int: malware_length, Int: seed, Int:  $d$   
 $ldpc \leftarrow \text{init\_ldpc}(\text{seed})$   
 $y \leftarrow []$   
 $\text{PNRG}(\text{seed})$   
 $\text{preamble} \leftarrow \text{random}([-1, 1], \text{size} = 200)$   
 $n \leftarrow \text{malware\_length}/d$   
 $i \leftarrow 0$   
 $j \leftarrow 0$   
while  $i < n$  do  
  while  $j < d$  do  
     $\text{code} \leftarrow \text{random}([-1, 1], \text{size} = \text{len}(W_i))$   
     $y_i \leftarrow \text{transpose}(\text{code}) * (W_i)$   
     $y.\text{append}(y_i)$   
     $j \leftarrow j + 1$   
  end  
   $i \leftarrow i + 1$   
end  
 $\text{gain} \leftarrow \text{mean}(\text{multiply}(y[:200], \text{preamble}))$   
 $\text{sigma} \leftarrow \text{std}(\text{multiply}(y[:200], \text{preamble})/\text{gain})$   
 $\text{snr} \leftarrow -20 * \log_{10}(\text{sigma})$   
 $\text{message} \leftarrow \text{ldpc.decode}(y[200:]/\text{gain}, \text{snr})$   
 $\text{malware} \leftarrow \text{message}[0 : \text{malware\_length}]$   
 $\text{hash} \leftarrow \text{message}[\text{malware\_length} :]$ 
```




$$W^{\text{Malicious}} = W + \gamma C b$$

γ — gain factor to control
the power of the signal

How does it perform?

— — —

Stealthiness



**Impact on
model
performance**




Robustness



Stealthiness

Stealthiness



**Impact on
model
performance**

Robustness

Stealthiness



1. Evaluation against anti-malware software.
2. Statistical analysis



Stealthiness



1. Evaluation against anti-malware software.

2. Statistical analysis

	Malware samples				
	Stuxnet	Destover	Bladabindi	Zeus	Kovter
Plain malware	89.19%	83.78%	72.97%	91.89%	62.16%
Stegomalware	0.00%	13.51%	8.11%	10.81%	5.41%
Malicious-DNN	0.00%	0.00%	0.00%	0.00%	0.00%


Detection rate on **Metadefender** metascan feature

Stealthiness



1. Evaluation against anti-malware software.

2. Statistical analysis




ResNet-101 architecture
Cerber malware

Comparison between the distribution of the weight parameters of the baseline regular model and the **one containing the malware**.

Stealthiness

— — —

Stealthiness




**Impact on
model
performance**

Robustness

Impact on model performance

— — —

Stealthiness



**Impact on
model
performance**

Robustness

Impact on model performance


	DenseNet		ResNet50		ResNet101		VGG11		VGG16	
Malware	Bas.	Mal.	Bas.	Mal.	Bas.	Mal.	Bas.	Mal.	Bas.	Mal.
Stuxnet	62.13	61.22	75.69	75.34	76.96	76.87	70.13	70.09	73.37	73.34
Destover	62.13	52.36	75.69	74.89	76.96	76.79	70.13	70.05	73.37	73.28

Classification accuracy (%) on the ImageNet task before and after embedding different sized malware into different sized neural networks.

Impact on model performance

— — —

Stealthiness




**Impact on
model
performance**

Robustness

Robustness

Stealthiness



**Impact on
model
performance**




Robustness




Robustness

1. Robustness towards fine-tuning
2. Robustness towards parameter pruning



Robustness

1. **Robustness towards fine-tuning**
2. Robustness towards parameter pruning



The change in signal-to-noise ratio of the malicious payload when the VGG11 architecture is trained on CIFAR10 and repurposed to solve the CIFAR100 task.

Robustness

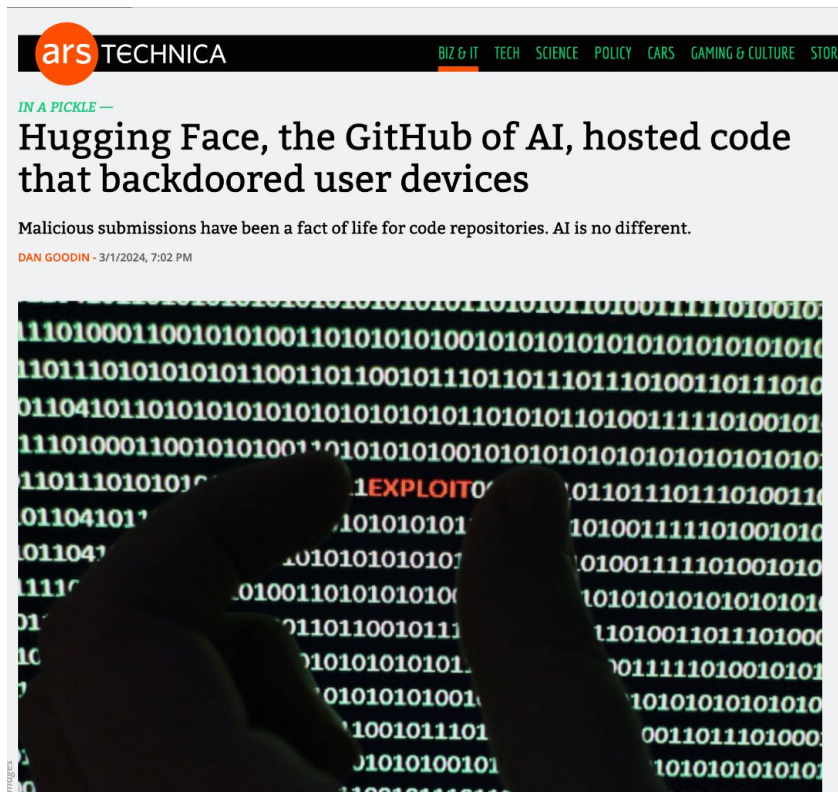
1. Robustness towards fine-tuning
2. **Robustness towards parameter pruning**

Pruning Ratio	Does the payload survive?
25.00%	YES
50.00%	YES
75.00%	YES
90.00%	NO
99.00%	NO

Fresh out of the oven (bad) news...



Hugging Face



<https://arstechnica.com/security/2024/03/hugging-face-the-github-of-ai-hosted-code-that-backdoored-user-devices/>



 Subscribe – Get Latest News

Contact



Feb 08, 2025 Ravie Lakshmanan

Artificial Intelligence / Supply Chain Security



Hugging Face

"The pickle files extracted from the mentioned PyTorch archives revealed the malicious Python content at the beginning of the file," ReversingLabs researcher Karlo Zanki [said](#) in a report shared with The

<https://thehackernews.com/2025/02/malicious-ml-models-found-on-hugging.html?m=1>

Fresh out of the oven (bad) news...

A Disney Worker Downloaded an AI Tool. It Led to a Hack That Ruined His Life.

Matthew Van Andel's experience reveals the threat that opportunistic hackers pose to companies and individuals

1 Matthew 'Dutch' Van Andel lost his job at Disney.

By [Robert McMillan](#) [Follow](#) and [Sarah Krouse](#) [Follow](#) | Photographs by Adali Schell for WSJ

Feb. 26, 2025 5:30 am ET

[Share](#) [AA](#) [Resize](#) [752](#)

[Listen](#) (2 min)

The stranger messaging Matthew Van Andel online last July knew a lot about him—including details about his lunch with co-workers at [Disney](#) [DIS 0.08%](#) [▲](#) from a few days earlier.

His mind raced; he knew no one outside Disney would have access to that

THE WALL STREET JOURNAL.

Most Popular News

What Went Wrong at Saudi Arabia's Futuristic Desert Metropolis



The Retirement-Savings Weapon Doctors and Lawyers Use to Build Wealth



Markets Finally Woke Up to Tariff Reality. Is This a Buying Opportunity?



Your Phone Needs a Burner



<https://www.wsj.com/tech/cybersecurity/disney-employee-ai-tool-hacker-cyberattack-3700c931>

Why so robust?



Binary sequence: `[0, 1, 1]`



PSK

`[-1, 1, 1]`

Spreading code: `[-1, 1, -1, -1, 1]`

Chip sequence: `[1, -1, 1, 1, -1, -1, 1, -1, -1, 1, -1, 1, -1, -1, 1]`

-5

+5

+5



Reading Material

1. Introduction to [Deep Neural Networks](#). (book chapter)
2. Digital communication and CDMA: [Link-1](#), [Link-2](#) (research papers).
3. Hiding malware into Deep Neural Networks [Link-1](#) [Link-2](#) (research papers).