

Security in Software Applications

Buffer Overflow



SAPIENZA
UNIVERSITÀ DI ROMA

Daniele Friolo

friolo@di.uniroma1.it

<https://danielefriolo.github.io>

Essence of the problem

Suppose in a C program we have an array of length 4:

```
char buffer[4]
```

What happens if we execute the statement below ?

```
buffer[4]='a'
```

Anything can happen !

If the data written (ie. the “a”) is user input that can be controlled by an attacker, this vulnerability can be exploited: anything that the attacker wants can happen.

Solution to the problem

Check array bounds at runtime

– Algol 60 proposed this back in 1960!

Unfortunately, C and C++ have not adopted this solution, for efficiency reasons.

(Ada, Perl, Python, Java, C#, and even Visual Basic have.)

As a result, buffer overflows have been the no 1 security problem in software ever since.

Problems caused by buffer overflows

Problems caused by buffer overflows:

- The first Internet worm, and all subsequent ones (CodeRed, Blaster, ...), exploited buffer overflows
- Buffer overflows cause in the order of 50% of all security alerts
 - Eg check out *CERT*, cve.mitre.org, or *bugtraq*
- Trends
 - Attacks are getting cleverer, defeating even more clever countermeasures
- Attacks are getting easier to do, by script kiddies

Problems caused by buffer overflows

Any C(++) code acting on untrusted input is at risk

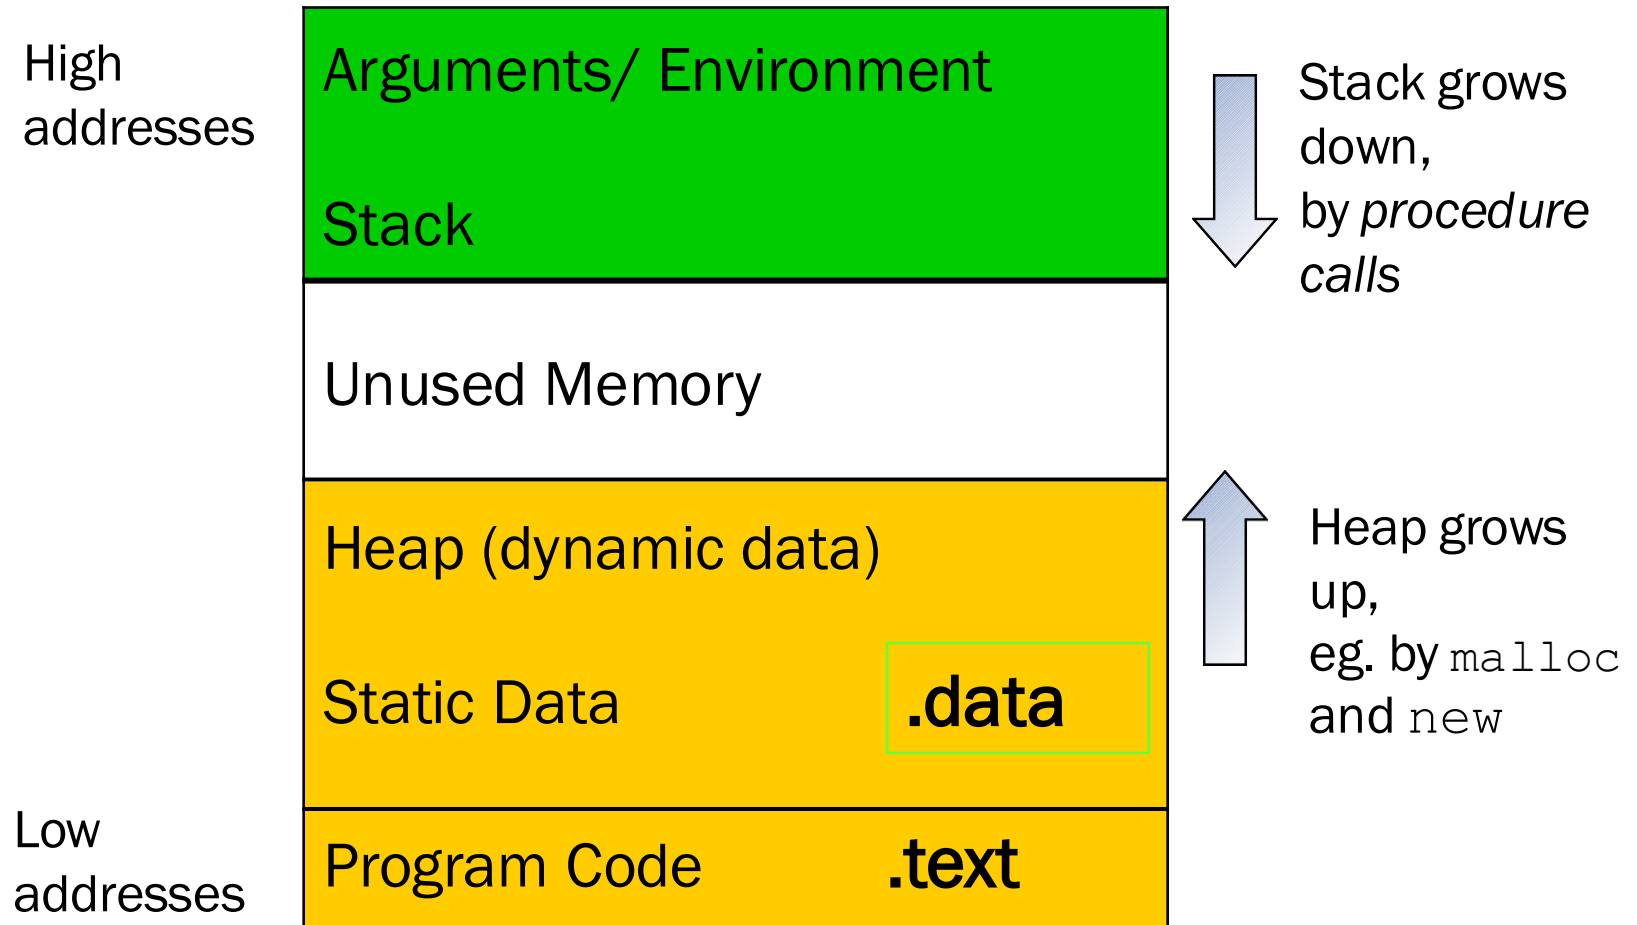
Eg

- code taking input over untrusted network
 - eg. sendmail, web browser, wireless network driver,...
- code taking input from untrusted user on multi-user system,
 - esp. services running with high privileges (as ROOT on Unix/Linux, as SYSTEM on Windows)
- code acting on untrusted files
 - that have been downloaded or emailed
- also embedded software, eg. in devices with (wireless) network connection such as mobile phones with Bluetooth, wireless smartcards, airplane navigation

Memory Management in C/C++

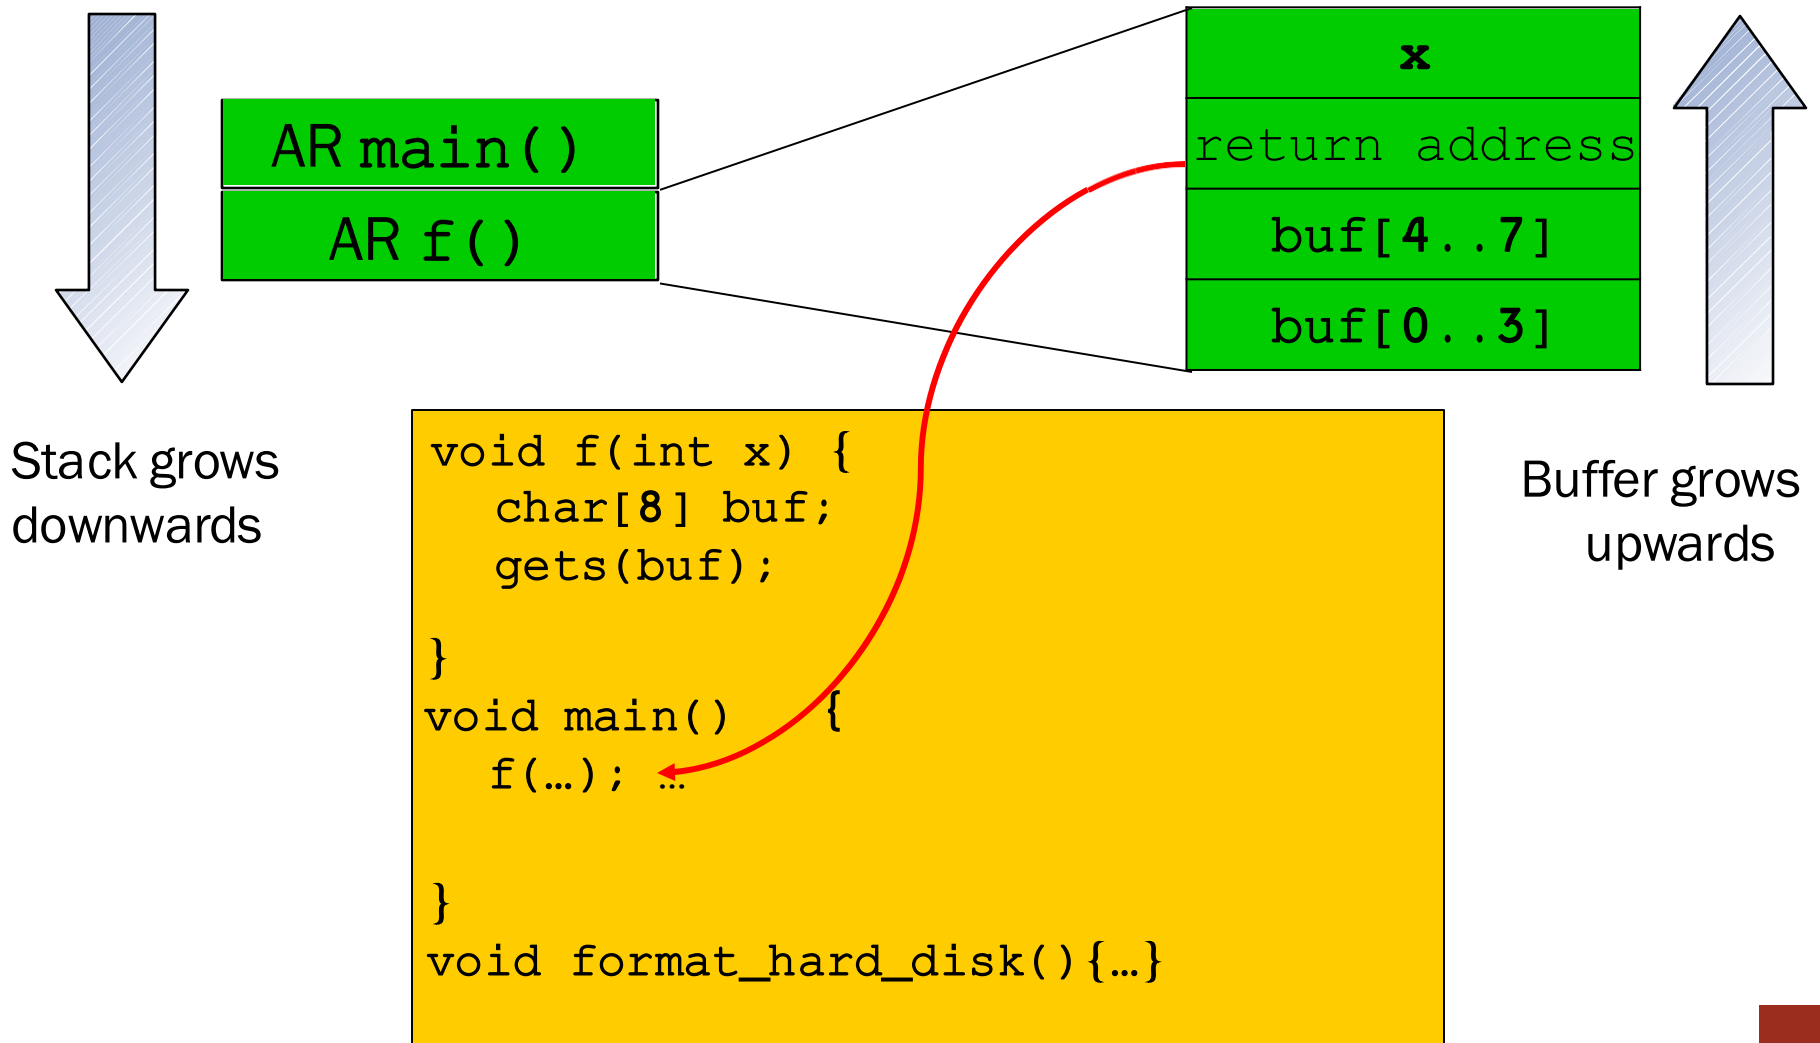
- Program responsible for its memory management
- Memory management is very error-prone
 - Who here has had a C(++) program crash with a segmentation fault?
- Typical bugs:
 - Writing past the bound of an array
 - Dangling pointers
 - missing initialisation, bad pointer arithmetic, incorrect de-allocation, double de-allocation, failed allocation, ...
- Memory leaks
- For efficiency, these bugs are not detected at runtime, as discussed before:
 - behaviour of a buggy program is undefined

Process Memory Layout



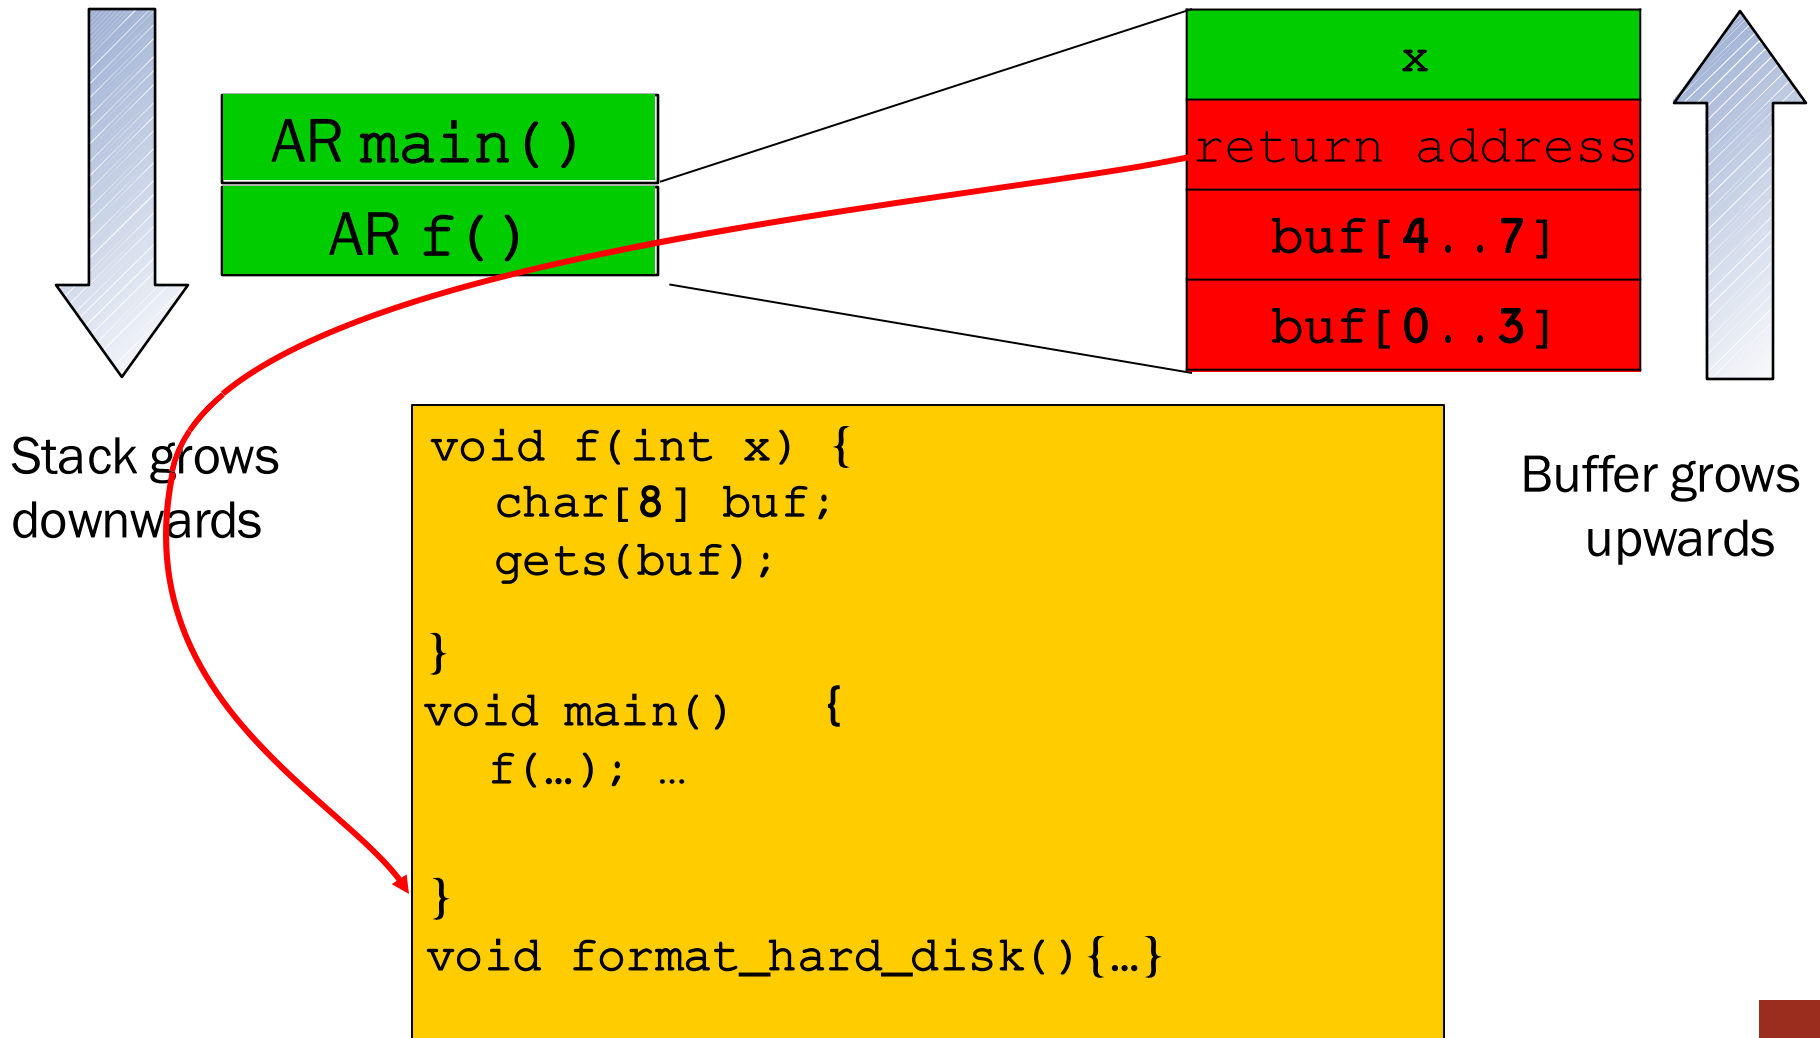
Stack Overflow

The stack consists of Activation Records:



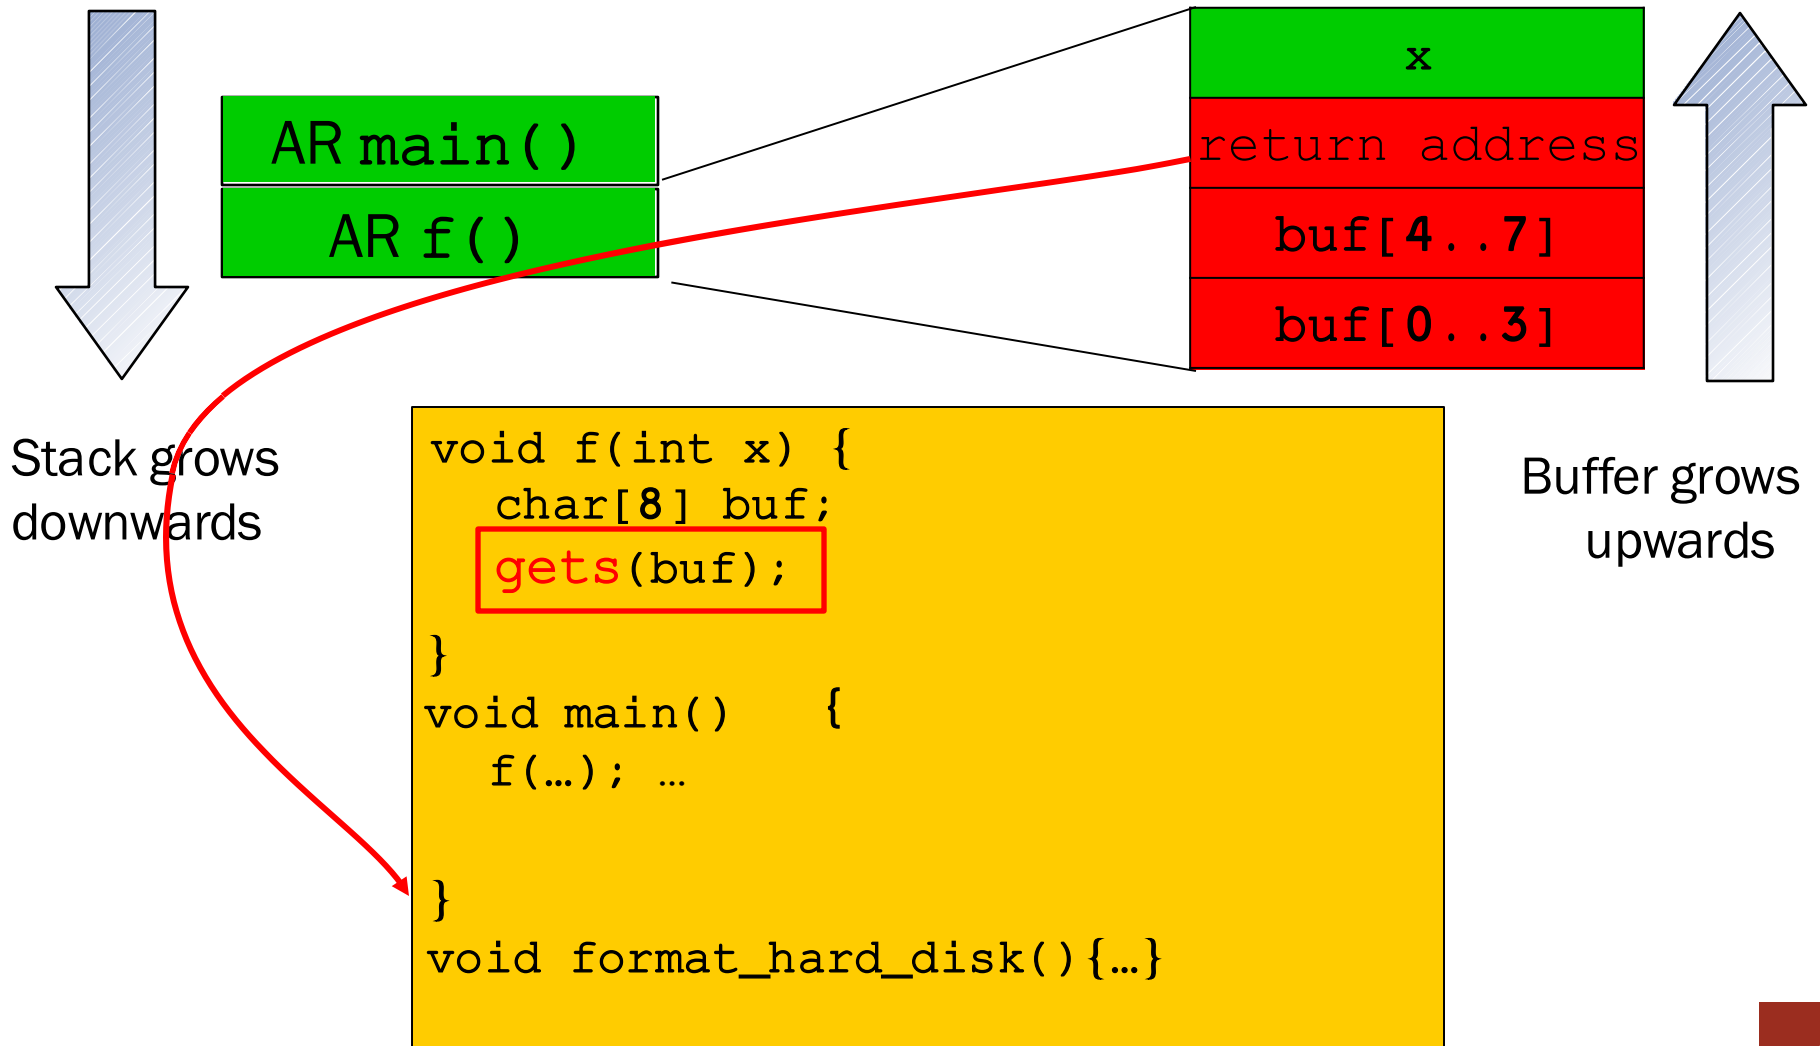
Stack Overflow

The stack consists of Activation Records:



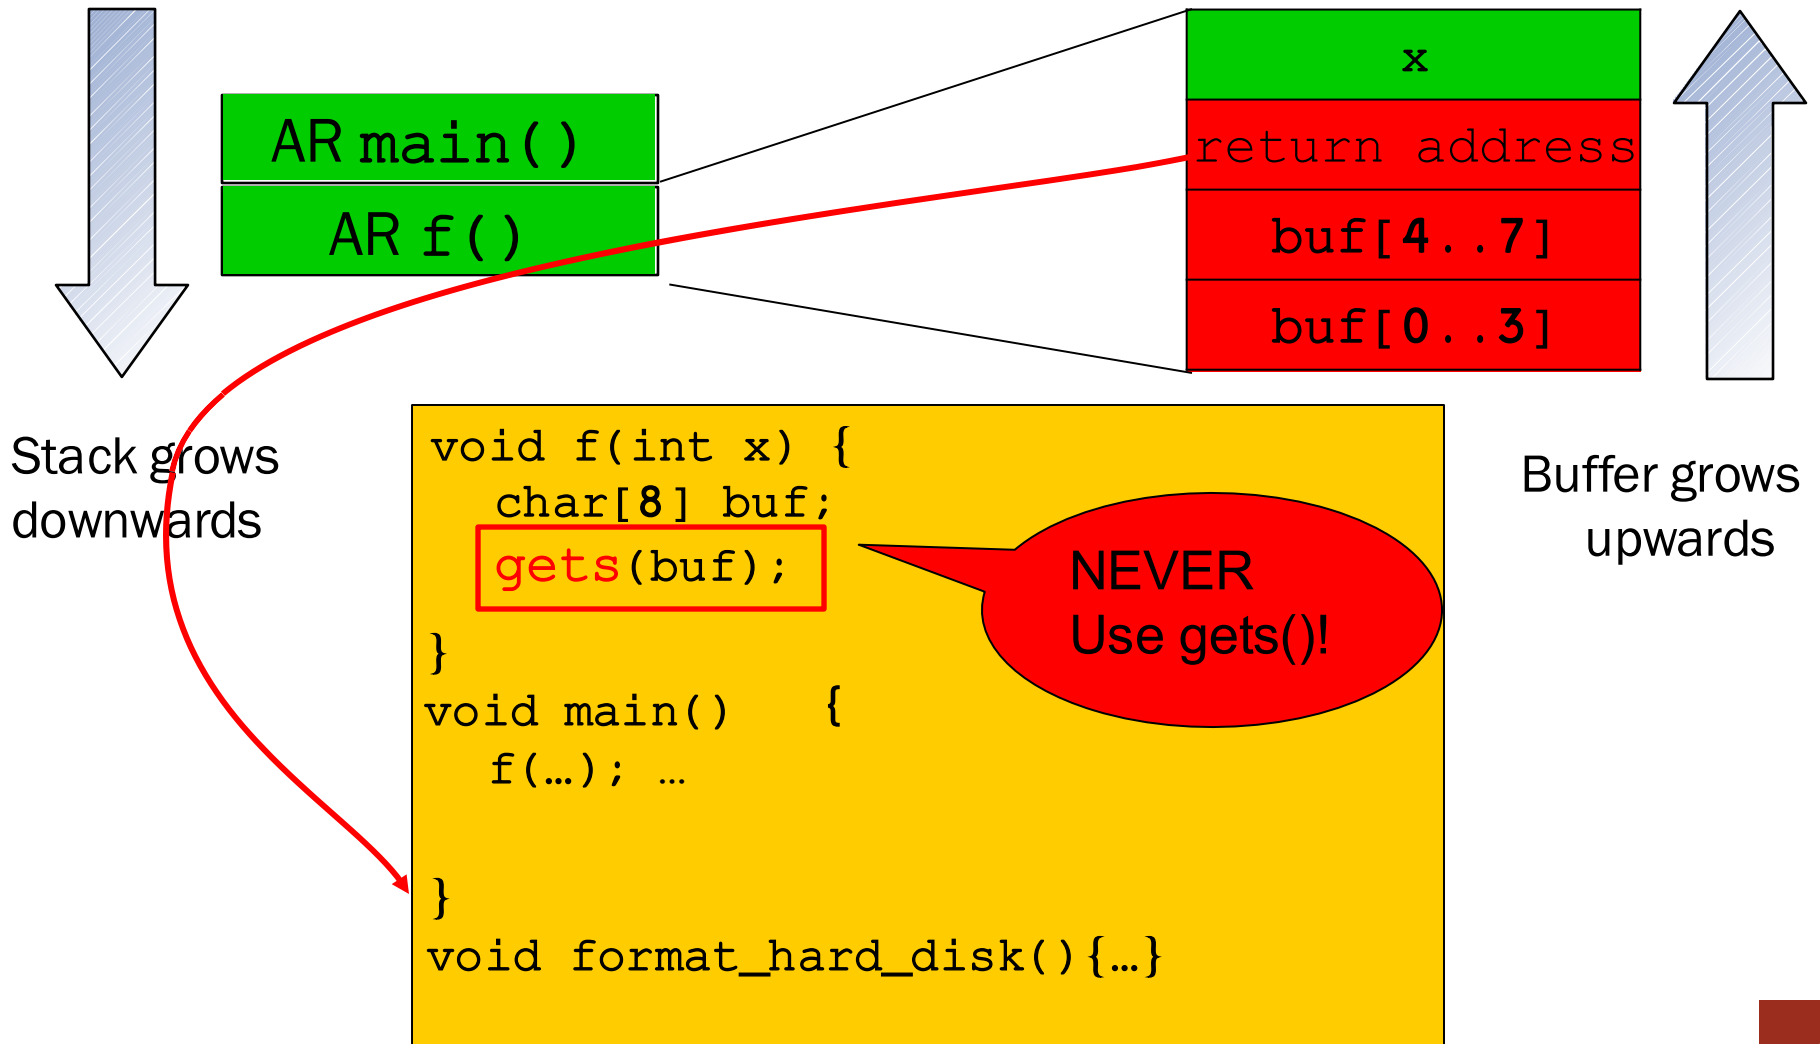
Stack Overflow

The stack consists of Activation Records:



Stack Overflow

The stack consists of Activation Records:



Stack Overflow

- Lots of details to get right:
 - No nulls in (character-)strings
 - Filling in the correct return address:
 - Fake return address must be precisely positioned
 - Attacker might not know the address of his own string
 - Other overwritten data must not be used before return from function
- ...

Stack Overflow

- **Stack overflow** is overflow of a buffer allocated on the stack
- **Heap overflow** idem, of buffer allocated on the heap

Common causes:

- poor programming with of *arrays* and *strings*
 - esp. library functions for null-terminated strings
- problems with *format strings*

But other low-level coding defects than can result in buffer overflows, eg
integer overflows or *data races*

Example: `gets()`

```
char buf[20];  
gets(buf); // read user input until  
           // first EoL or EOF character
```

- *Never* use `gets`
- Use `fgets(buf, size, stdin)` instead

Example: strcpy()

```
char dest[20];  
strcpy(dest, src); // copies string src to dest
```


- **strcpy** assumes that
 - **dest** is long enough
 - **src** is null-terminated
- Use **strncpy(dest, src, size)** instead

Spot the defect

```
char buf[20];  
char prefix[] = "http://";  
...  
strcpy(buf, prefix);  
    // copies the string prefix to buf  
strncat(buf, path, sizeof(buf));  
    // concatenates path to the string buf
```


Spot the defect

```
char buf[20];  
char prefix[] = "http://";  
...  
strcpy(buf, prefix);  
    // copies the string prefix to buf  
strncat(buf, path, sizeof(buf));  
    // concatenates path to the string buf
```



strncat's 3rd parameter is number of chars to copy, not the buffer size

Another common mistake is giving
sizeof(path) as 3rd argument...

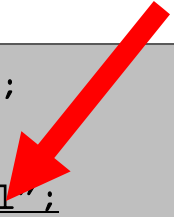
Spot the defect

```
char src[9]; char dest[9];

char base_url = "www.ru.nl";
strncpy(src, base_url, 9);
    // copies base_url to src
strcpy(dest, src);
    // copies src to dest
```

Spot the defect

base_url is 10 chars long, incl. its null terminator,
so **src** won't be null-terminated



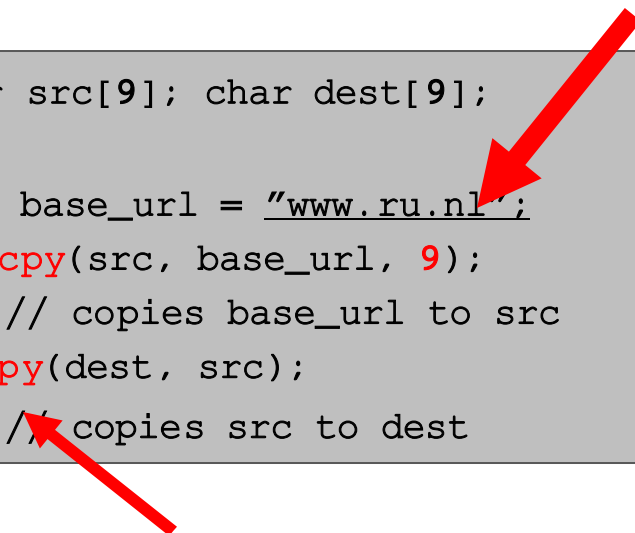
```
char src[9]; char dest[9];

char base_url = "www.ru.nl";
strncpy(src, base_url, 9);
    // copies base_url to src
strcpy(dest, src);
    // copies src to dest
```

Spot the defect

base_url is 10 chars long, incl. its null terminator,
so **src** won't be null-terminated

```
char src[9]; char dest[9];  
  
char base_url = "www.ru.nl";  
strncpy(src, base_url, 9);  
    // copies base_url to src  
strcpy(dest, src);  
    // copies src to dest
```



so **strcpy** will overrun the
buffer **dest**

Example: strcpy and strncpy

- Don't replace

```
strcpy(dest, src)
```

by

```
strncpy(dest, src, sizeof(dest))
```

but by

```
strncpy(dest, src, sizeof(dest) - 1)
```

```
dst[sizeof(dest) - 1] = `\\0`;
```

if dest should be null-terminated!

- Btw: a *strongly typed programming language* could of course enforce that strings are always null-terminated...

Spot the defect

```
char *buf;
int i, len;

read(fd, &len, sizeof(len));
    // read sizeof(len) bytes, ie. an int
    // and store these in len
buf = malloc(len);
read(fd, buf, len);
```

Spot the defect

We forget to check for bytes representing a negative int, so **len** might be negative

```
char *buf;
int i, len;

read(fd, &len, sizeof(len));
    // read sizeof(len) bytes, ie. an int
    // and store these in len
buf = malloc(len);
read(fd, buf, len);
```

len cast to unsigned and negative length overflows

read then goes beyond the end of **buf**

Spot the defect

```
char *buf;
int i, len;

read(fd, &len, sizeof(len));
    // read sizeof(len) bytes, ie. an int
    // and store these in len
buf = malloc(len);
read(fd, buf, len);
```

Remaining problem may be that **buf** is not null-terminated

Spot the defect

```
char *buf;
int i, len;

read(fd, &len, sizeof(len));
    // read sizeof(len) bytes, ie. an int
    // and store these in len
buf = malloc(len);
read(fd, buf, len+1);
buf[len] = '\0'; // null terminate buf
```

Absence of language-level security

In programming languages with “security” provisions, the programmer would not have to worry about

- *writing past the bounds of the array* (`IndexOutOfBoundsException` for example)
- *implicit conversion from signed to unsigned integers* (forbidden or warned by compiler/typechecker)
- `malloc` returning null value (`OutOfMemoryException` for example)
- `malloc` non initializing memory (by default)
- *integer overflow* (`IntegerOverflowException` for example)

Spot the defect

```
#ifdef UNICODE
#define _sntprintf _snwprintf #define TCHAR
wchar_t
#else
#define _sntprintf _snprintf #define TCHAR
char
#endif

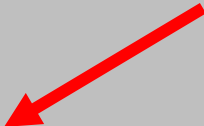
TCHAR buff[MAX_SIZE];
_sntprintf(buff, sizeof(buff), "%s\n", input);
```

Spot the defect

```
#ifdef UNICODE
#define _sntprintf _snwprintf #define TCHAR
wchar_t
#else
#define _sntprintf _snprintf #define TCHAR
char
#endif
```

*_snwprintf's 2nd param is # of chars in
buffer, not # of bytes*

```
TCHAR buff[MAX_SIZE];
_sntprintf(buff, sizeof(buff), "%s\n", input);
```



The *CodeRed* worm exploited such an ANSI/Unicode mismatch

Spot the defect

```
#define MAX_BUF = 256

void BadCode (char* input)
{
    short len;
    char buf[MAX_BUF];

    len = strlen(input);

    if (len < MAX_BUF) strcpy(buf,input);
}
```

Spot the defect

```
#define MAX_BUF = 256
```

```
void BadCode (char* input)
```

```
{    short len;
```

```
    char buf[MAX_BUF];
```

```
    len = strlen(input);
```

```
    if (len < MAX_BUF) strcpy(buf, input);
```


```
}
```

What if **input** is longer than 32K ?

len will be a negative number, due to integer overflow

hence: potential *buffer overflow*

The *integer overflow* is the root problem, but the (*heap*) *buffer overflow* that this enables make it exploitable

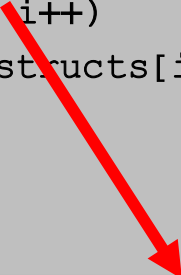


Spot the defect

```
bool CopyStructs(InputFile* f, long count)
{
    structs = new Structs[count];
    for (long i = 0; i < count; i++)
        { if !(ReadFromFile(f,&structs[i])) )
            break;
        }
}
```

Spot the defect

```
bool CopyStructs(InputFile* f, long count)
{
    structs = new Structs[count];
    for (long i = 0; i < count; i++)
        { if !(ReadFromFile(f,&structs[i])) )
            break;
        }
}
```



effectively does a
`malloc(count*sizeof(type))` which
may cause *integer overflow*

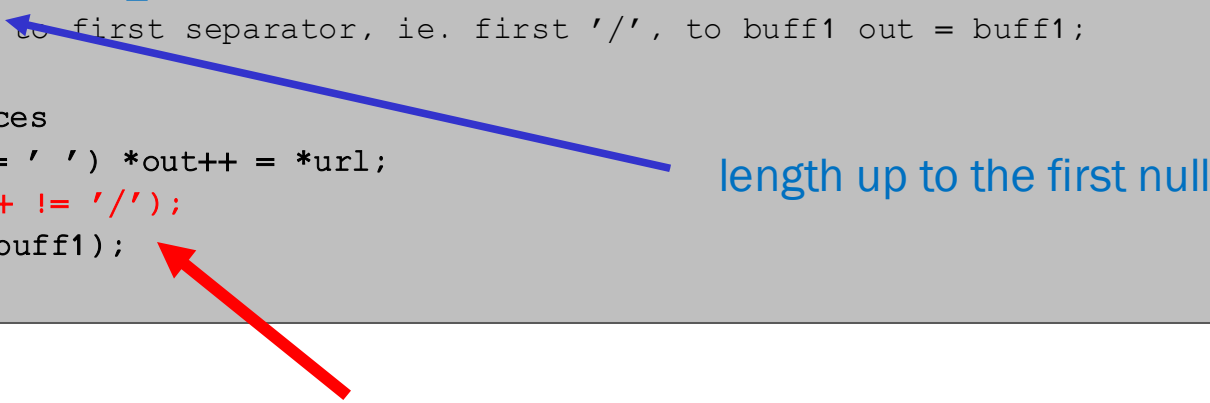
And this integer overflow can lead to a (heap) buffer overflow.
(Microsoft Visual Studio 2005(!) C++ compiler adds check to prevent this)

Spot the defect

```
char buff1[MAX_SIZE], buff2[MAX_SIZE];
// make sure url a valid URL and fits in buff1 and buff2:
if (! isValid(url)) return;
if (strlen(url) > MAX_SIZE - 1) return;
// copy url up to first separator, ie. first '/', to buff1 out = buff1;
do {
    // skip spaces
    if (*url != ' ') *out++ = *url;
} while (*url++ != '/');
strcpy(buff2, buff1);
...
```

Spot the defect

```
char buff1[MAX_SIZE], buff2[MAX_SIZE];
// make sure url a valid URL and fits in buff1 and buff2:
if (! isValid(url)) return;
if (strlen(url) > MAX_SIZE - 1) return;
// copy url up to first separator, ie. first '/', to buff1 out = buff1;
do {
    // skip spaces
    if (*url != ' ') *out++ = *url;
} while (*url++ != '/');
strcpy(buff2, buff1);
...
```



length up to the first null

what if there is no '/' in the URL?

Spot the defect

```
char buff1[MAX_SIZE], buff2[MAX_SIZE];
// make sure url a valid URL and fits in buff1 and buff2:
if (! isValid(url)) return;
if (strlen(url) > MAX_SIZE - 1) return;
// copy url up to first separator, ie. first '/', to buff1 out = buff1;
do {
    // skip spaces
    if (*url != ' ') *out++ = *url;
} while (*url++ != '/')
&& (*url != 0);
strcpy(buff2, buff1);
...
```

Spot the defect

```
char buff1[MAX_SIZE], buff2[MAX_SIZE];
// make sure url a valid URL and fits in buff1 and buff2:
if (! isValid(url)) return;
if (strlen(url) > MAX_SIZE - 1) return;
// copy url up to first separator, ie. first '/', to buff1 out = buff1;
do {
    // skip spaces
    if (*url != ' ') *out++ = *url;
} while (*url++ != '/')
&& (*url != 0);
strcpy(buff2, buff1);
...
```

Order of tests is wrong (note the first test includes ++)

What about 0-length URLs?

Is buff1 always null-terminated?

Spot the defect

```
#include <stdio.h>

int main(int argc, char* argv[])
{
    if (argc > 1)
        printf(argv[1]);
    return 0;
}
```

This program is vulnerable to *format string attacks*, where calling the program with strings containing special characters can result in a buffer overflow attack.

Format String attacks

- Complete new type of attack, invented/discovered in 2000. Like integer overflows, it can lead to buffer overflows.
- Strings can contain special characters, eg %s in

```
printf("Cannot find file %s", filename);
```

Such strings are called format strings
- What happens if we execute the code below?

```
printf("Cannot find file %s");
```
- What *may* happen if we execute

```
printf(string)
```

where **string** is user-supplied ?
Esp. if it contains special characters, eg %s , %x , %n , %hn?

Format String attacks

- `%x` reads and prints 4 bytes from stack
 - this may leak sensitive data
- `%n` writes the number of characters printed so far onto the stack
 - this allow stack overflow attacks...
- Note that format strings break the “don’t mix data & code” principle.
- Easy to spot & fix:
replace `printf(str)` **by** `printf("%s", str)`

Dynamic Countermeasures

Protection by kernel

- *Non-executable memory* (NOEXEC)
 - Prevents attacker executing her code
- *Address space layout randomisation* (SLR)
 - Generally makes attacker's life harder
- *Instruction set randomisation*
 - Hardware support needed to make this efficient enough

Protection inserted by the compiler

- *Stack canaries* to prevent or detect malicious changes to the stack; examples to follow
- *Obfuscation* of memory addresses

Doesn't prevent against heap overflows

Dynamic Countermeasures: Stack Canaries

- introduced in *StackGuard* in gcc
- a dummy value - *stack canary or cookie* - is written on the stack in front of the *return address* and checked when function returns
- a careless stack overflow will overwrite the canary, which can then be detected.
- a careful attacker can overwrite the canary with the correct value.
- additional countermeasures:
 - use a random value for the canary
 - XOR this random value with the return address
 - include string termination characters in the canary value

Further Improvements

- ***PointGuard***
 - also protects other data values, eg function pointers, with canaries
- ***ProPolice's Stack Smashing Protection (SSP)*** by IBM
 - also *re-orders stack elements* to reduce potential for trouble
- ***Stackshield*** has a special stack for return addresses, and can disallow function pointers to the data segment

Dynamic Countermeasures

NB none of these protections is perfect!

Eg

- even if attacks to return addresses are caught, integrity of other data other the stack can still be abused
- clever attacks may leave canaries intact
- where do you store the "master" canary value
 - a cleverer attack could change it
- none of this protects against heap overflows
 - eg buffer overflow within a struct...
-

Other Countermeasures

- We can take countermeasures at different points in time
 - before we even begin programming
 - during development
 - when testing
 - when executing code
- to prevent, to detect – at (pre)compile time or at runtime -, and to mitigate problems with buffer overflows

Prevention

- Don't use C or C++
- Better programmer awareness & training

Eg read – and make other people read -

- Building Secure Software, J. Viega & G. McGraw, 2002
- Writing Secure Code, M. Howard & D. LeBlanc, 2002
- 19 deadly sins of software security, M. Howard, D LeBlanc & J. Viega, 2005
- Secure programming for Linux and UNIX HOWTO, D. Wheeler,
- Secure C coding, T. Sirainen
- The Secure Coding Cookbook for C and C++ by John Viega and Matt Messier
- Secure Coding: Principles and Practices by Robert Seacord, 2013

Dangerous C system calls

Extreme risk

- gets

High risk

- strcpy
- strcat
- sprintf
- scanf
- sscanf
- fscanf
- vfscanf
- vsscanf

- streadd
- strcpy
- strtrns
- realpath
- syslog
- getenv
- getopt
- getopt_long
- getpass

Moderate risk

- getchar
- fgetc
- getc
- read
- bcopy

Low risk

- fgets
- memcpy
- snprintf
- strncpy
- strccpy
- strcadd
- strncpy
- strncat
- vsnprintf

Prevention – Use better libraries

- there is a choice between using statically vs dynamically allocated buffers
 - *static* approach easy to get wrong, and chopping user input may still have unwanted effects
 - *dynamic* approach susceptible to out-of-memory errors, and need for failing safely

Prevention – Use better libraries (strings)

- `libsafe.h` provides safer, modified versions of eg `strcpy`
 - prevents buffer overruns beyond current stack frame in the dangerous functions it redefines
- `libverify` enhancement of `libsafe`
 - keeps copies of the stack return address on the heap, and checks if these match

`strncpy(dst,src,size)` and `strncat(dst,src,size)`

with `size` the size of `dst`, not the maximum length copied.

Consistently used in OpenBSD

Prevention – Use better libraries (strings)

- `glib.h` provides `Gstring` type for dynamically growing null-terminated strings in C
 - but failure to allocate will result in crash that cannot be intercepted, which may not be acceptable

`Strsafe.h` by Microsoft guarantees null-termination and always takes destination size as argument

C++ `string` class

- but `data()` and `c_str()` return low level C strings, ie `char*`, with result of `data()` is not always null-terminated on all platforms...

Detection (before shipping)

- Testing
 - Difficult! How to hit the right cases?
 - *Fuzz testing* - test for crash on long, random inputs – can be succesful in finding some weaknesses
- Code reviews
 - Expensive & labour intensive
- Code scanning tools (aka static analysis) Eg
 - *RATS* – also for PHP, Python, Perl
 - *Flawfinder* , *ITS4*, *Deputy*, *Splint*
 - *PREfix*, *PREfast* by Microsoft plus other commercial tools

More prevention & detection

The most extreme form of static analysis:

- Program verification
 - Proving by mathematical means (eg hoare logic) that memory management of a program is safe
 - Extremely labour-intensive
 - E.g. hypervisor verification project by microsoft & verisoft.

<https://www.Microsoft.Com/en-us/research/project/vcc-a-verifier-for-concurrent-c/>

Reducing attack surface

- Not running or even installing certain software, or enabling all features by default, mitigates the threat

Summary

- Buffer overflows are the top security vulnerability
- Any C(++) code acting on untrusted input is at risk
- Getting rid of buffer overflow weaknesses in C(++) code is hard (and may prove to be impossible)
 - Ongoing arms race between countermeasures and ever more clever attacks.
 - Attacks are not only getting cleverer, using them is getting easier

More general

Buffer overflow is an instance of three more general problems:

- 1) Lack of *input validation*
- 2) Mixing *data & code*
 - Data and return address on the stack
- 3) Believing in & relying on an *abstraction*
 - In this case, the abstraction of procedure calls offered by C

Attacks often exploit holes in abstractions that are not 100% enforced

Moral of the story

- Don't use C(++), if you can avoid it
 - But use a language that provides memory safety, such as java or C#
- If you do have to use c(++), become or hire an expert
- **Reading**
 - *A Comparison of Publicly Available Tools for Dynamic Buffer Overflow Prevention*, by John Wilander and Mariam Kakkar